Nova AI -- Product Documentation

Nova Interview Architecture

Nova Interview is built on a microservices architecture with the following core components: the Interview Orchestrator (LangGraph-based state machine), the Voice Service (WebSocket real-time audio with ElevenLabs TTS and Whisper STT), the Evaluation Engine (LLM-graded scoring across technical and behavioral dimensions), and the Knowledge Base (RAG pipeline with Pinecone vector store for company-specific context).

Key Features

Multi-mode Interviews:

Candidates can choose between voice-only, chat-only, or mixed-mode interviews. Voice interviews use WebSocket streaming with sub-500ms latency.

Adaptive Questioning:

The engine dynamically adjusts question difficulty based on candidate performance. Strong answers trigger deeper follow-ups; weak answers trigger scaffolding hints.

Structured Scorecards:

Every interview produces a standardized scorecard across 6 competencies: problem-solving, technical knowledge, communication, system design, coding, and culture fit.

Batch Processing:

Enterprise plans support concurrent interview sessions with automatic load balancing across LLM providers (OpenRouter, Groq, OpenAI).

Integration API

Nova AI exposes RESTful APIs for ATS integration (Greenhouse, Lever, Workday), single sign-on (Google OAuth, Microsoft Entra), and webhook-based event notification. The API supports CSV/JSON export of interview results and scorecard data.

Security & Compliance

- SOC 2 Type II certified (audit completed February 2026)
- GDPR compliant with data residency in US, EU, and APAC regions
- All candidate data encrypted at rest (AES-256) and in transit (TLS 1.3)
- Role-based access control with audit logging for all administrative actions